PPCWay

Research report on Otto AI

Build Otto on Claude, then settle residency

Which language model should power Otto, PPCWay's Google Ads assistant, and is there one built for Google Ads?

Prepared for	The PPCWay owner and the client
Date	14 October 2026
Checked	Fact-checked against live sources the same day

Contents

At a glance

1 Otto's accuracy is decided by its tools more than its model

2 Claude Sonnet 5.5 leads a three-model lineup at about $12 a customer

3 No purpose-built Google Ads model beats a general model with good tools

4 Only one ageing Gemini model keeps Otto's processing inside Canada

5 The client's decision is what "stored in Canada" has to mean

Conclusion

What the fact-check changed

Sources

At a glance

Run Otto on Anthropic's Claude models, called through Google Cloud's Vertex AI in its US region, and keep customer data stored in Canada. On a typical small-business workload this costs about $11.60 per customer a month, under 8% of the $149 Starter plan.

Otto's job	Model	Per customer a month
Answers questions about the account	Claude Sonnet 5.5	$8.14
Explains the daily recommendations, overnight	Claude Sonnet 5.5	$1.80
Writes ad text and keyword drafts	Claude Opus 5.5	$0.60
Routes requests and runs small internal jobs	Claude Haiku 5.5	$0.04
Google Cloud's US-region charge, 10%		$1.06
Total		**$11.64**

Which model should power Otto?

Claude Sonnet 5.5 for the chat, Haiku 5.5 for small jobs and Opus 5.5 for ad text. OpenAI's GPT-6.1 Sol is the close alternative at the same price. A one- to two-day test on about 150 real Otto questions settles it.

Is there a model built for Google Ads?

No usable one. Google's own Ads assistant has no API, and every PPC competitor examined runs on general models from OpenAI, Anthropic or Google. Otto's accuracy comes from the design around the model: read-only tools over the Postgres mirror, every number worked out in code, and a check on each number in a reply. Two Google tools are worth a look alongside Claude: the Ads API's ad-text generator, if PPCWay gets into its closed beta, and Google's coding assistant for the engineers.

What must the client decide?

Whether "data stored in Canada" also covers where Otto's AI runs. No current frontier model processes data in Canada. A Canada-only Otto would run on Gemini 3.5 Flash with reserved capacity, at about $23–31 per customer, and Google has not confirmed it sells that capacity in Montréal. The report recommends storing data in Canada, processing in the US, and rewording the promise.

What happens next?

1. Run the comparison test on Sonnet 5.5, GPT-6.1 Sol and Gemini 3.5 Flash, against a frozen copy of the Google Ads test accounts.
2. The client chooses what "stored in Canada" covers, from the three options in section 5.
3. Before launch, complete the six items in section 5, from blocking Google's global endpoint to updating the privacy notice.

1

Otto's accuracy is decided by its tools more than its model

The biggest factor in whether Otto tells the truth about an account is how the data reaches the model, not which model reads it.

In dbt Labs' April 2026 re-test (11 questions, each run 20 times), Claude Sonnet 4.6 scored **90.0% writing its own SQL against 98.2% when it asked for defined metrics through a semantic layer**, and GPT-5.3-Codex went from 84.1% to 100%. On questions the semantic layer did not cover, it returned an error rather than a wrong number. With text-to-SQL, in dbt's words, "failure looks like a plausible but incorrect answer"[1]. An independent paired study of 99 questions found the same pattern. Adding a short hand-written markdown description of what the metrics mean lifted accuracy from 45.5–50.5% to 67.7–68.7%, and within each condition Claude Opus 4.7, Claude Sonnet 4.6 and GPT-5.4 were statistically indistinguishable: "model choice within tier does not" explain the variance[2]. Even the best text-to-SQL system on the BIRD benchmark reaches **82.95% against 92.96% for humans**[3]. For a customer who cannot read SQL, that is roughly one wrong answer in six.

Otto should therefore work through a small set of typed, read-only tools over the Postgres mirror, not write queries itself. Example tools are `get_performance(entity, date_range, metrics)`, `compare_periods`, `top_search_terms(filter="wasted")` and `explain_change`. Shopify found that up to about 20 tools stay clear and easy to debug, that 20–50 tools blur together, and that more than 50 produce several ways to do the same thing[4]. Each tool returns figures already computed and formatted in code, each with an id. A question like "why did my cost per lead rise?" goes to a fixed calculation that splits the change into its main drivers and returns the top three. The model only puts those drivers into words. Tableau Pulse uses the same approach after finding that raw metric facts fed to a model "would often be modified incorrectly, entirely omitted, or included in a repetitive manner"[5]. A final check in the FastAPI layer pulls every number out of the draft reply and matches it against that turn's tool results. If one does not match, Otto rewrites the reply once on Opus 5.5, then falls back to a template. The design behind this check, known as proof-carrying numbers, makes the point that retrieval, citations and uncertainty estimates "improve transparency but cannot guarantee fidelity"[6]. The approach also meets the requirement that every recommendation "names the data it used" (Requirements).

The same design keeps the rule that the model never calls Google Ads. Otto's only tool that leads to a change is `propose_change`. It writes a draft into the proposed-changes table and nothing more. Approval, the guardrails and the gateway do the rest. This also suits Google's own rules. The Google Ads API policies say "end users of your tool will need to manually sign in to make changes to their accounts", and prohibit "headless programmatic access to make account modifications" without Google's review of the use case or "direct, per-entity Google authentication". They also bar a hosted "MCP

server, proxy endpoint, or any similar service that solely replicates, wraps, or re-exposes Google Ads programmatic capabilities", unless Google has reviewed it. Google's review for such interfaces asks for "immutable transaction logs that map every automated programmatic action directly to a verified end-advertiser account"[7]. PPCWay's append-only action log already does that, which helps if PPCWay ever needs that review. The search-terms report needs extra care. Anyone can type instructions into Google, and if that search triggers a customer's ad, the text lands in the mirror and later in Otto's context. Prompt injection is still OWASP's number one risk for AI applications[8]. The agreed defence is that once an agent has read untrusted text, "it must be constrained so that it is impossible for that input to trigger any consequential actions"[9]. So tools mark synced text as untrusted and cut it short, Otto has no outbound channel such as fetching a web address, and the account comes from the signed-in session, enforced by Postgres row-level security, never from the model.

All of this changes what to look for in a model. Once the numbers are computed in code, the model mainly has to pick the right tool with the right arguments across several turns, follow the reply format (a sentence, then the number), start replying quickly and stay cheap. Raw reasoning power and long context matter little. In the dbt test, mid-tier Sonnet 4.6 beat flagship Opus 4.6 on semantic-layer questions. Raising GPT's reasoning effort from high to xhigh took answers from about 8 to over 20 seconds with no gain in accuracy, and for Claude, effort "changed neither results nor latency"[1].

2

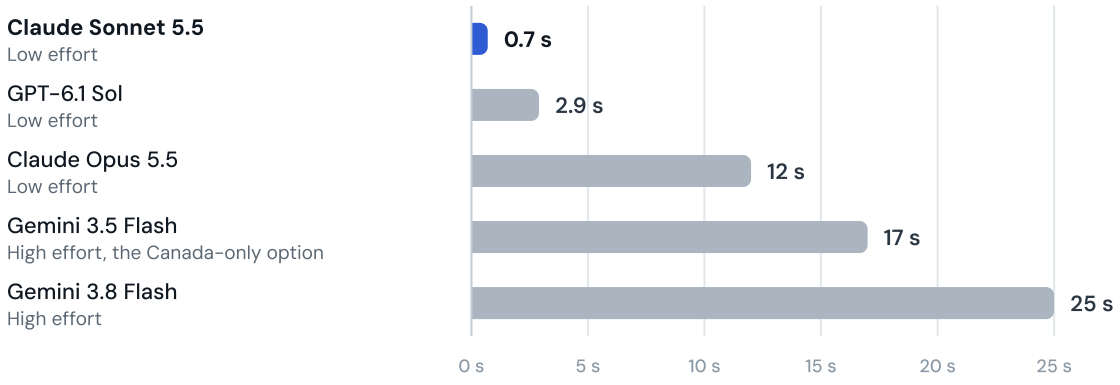
Claude Sonnet 5.5 leads a three-model lineup at about $12 a customer

Claude suits Otto best for three reasons: it is the quickest capable model to start answering, it led the last independent tool-calling tests, and the product is already built around it.

One caution applies to all the evidence that follows. Anthropic, OpenAI and Google each replaced their whole lineup between 2 September and 7 October 2026, so independent benchmarks for the current models are thin and some figures come from the vendors themselves[10].

Sonnet 5.5 costs **$2 per million input tokens and $10 per million output tokens**. Reading from the prompt cache costs $0.10 per million, cut from $0.20 on 7 October ($0.11 in Vertex's US region)[11,12]. On Artificial Analysis's live measurements on 14 October, Sonnet 5.5 at low effort began answering in **about 0.7 seconds**, against about 2.9 seconds for GPT-6.1 Sol at low effort and about 25 seconds for Gemini 3.8 Flash at high. These figures move from day to day[13]. The Berkeley function-calling leaderboard was last updated in April 2026 and only covers the previous generation. On it, Claude Opus 4.5 led at 77.47% overall and 68.38% on multi-turn tasks, while GPT-5.2 managed 55.87% overall and only 28.12% multi-turn[14]. On the GDPval-AA knowledge-work test that Anthropic reports, Sonnet 5.5 scored 1844 against 1487 for GPT-6 Sol, and Haiku 5.5 scored 1620 against 1437 for GPT-6 Luna[15,12]. Anthropic's own footnotes say the Sonnet run used a pre-release model with a structured-output bug and that the GPT-6 Sol score may predate an OpenAI fix, so treat these as direction, not proof. On fit, the Blueprint already chose Anthropic with Opus for every writing step (Blueprint §7.2). The sign-up consent text and the designed "companies that process your data" screen already name Anthropic (settings screen).

Seconds before the model starts answering



Artificial Analysis, measured 14 October 2026. These figures move from day to day; the order has held.

Otto job	Model	Why this one	Est. cost per customer-month
Chat about the account (30 conversations of 6 turns)	Sonnet 5.5, low or medium effort, cached prompt	Quickest first reply in the mid tier; thinking can be switched off	$8.14
Daily recommendation explanations	Sonnet 5.5 through Vertex AI batch prediction (50% off)	The rules engine decides and the model only explains	$1.80
Ad text, keyword grouping and other drafts	Opus 5.5	The Blueprint's "ad copy is the product"; now $4/$20 against Opus 5's $5/$25	$0.60
Routing, titles, trimming long conversations, English summaries	Haiku 5.5	$0.10/$0.50, so these jobs cost cents	$0.04
Rewrite when a reply fails the number check	Opus 5.5	Highest score on Artificial Analysis's knowledge-reliability test (46, at max effort)	rare
Total		**Global routing $10.58; US region (+10%) $11.64**	

The per-job figures come from a cost model with stated assumptions: 30 conversations a month of 6 turns, about 8,000 tokens of tool results per turn, the cached system prompt reused, and recommendations sent as an overnight batch[11]. The batch runs through Google's own batch prediction, which reads from BigQuery or Cloud Storage. Anthropic's Message Batches API is not offered on Google Cloud. Google lists Sonnet 5.5 batch at $1.10/$5.50 per million tokens in the US multi-region, half the on-demand rate. Its batch guide does not name `us` as a location outright, so one test job should confirm it[16,17,18]. Recommendation text moves off Opus because the Blueprint itself called moving the narration tasks off Opus (to Haiku) its lowest-risk saving. Sonnet is the more cautious version of that step (Blueprint §19.3). The Blueprint's "Opus for everything" default should not be extended to chat. Opus 5.5 takes **about 12 seconds to start a reply even at low effort**, and using it for every job doubles the bill[13]. For comparison, the Blueprint estimated $0.20–0.60 per merchant a month for the writing jobs alone (Blueprint §7.7). The chat is what raises the bill to about $11.60, because every question re-reads the account's figures.

Model setup	Per customer-month	Share of Starter $149	Share of Growth $349	Agency, 10 accounts (share of $949)
Recommended Claude setup, US region	**$11.64**	**7.8%**	**3.3%**	**$116 (12.3%)**
GPT-6.1 Sol with GPT-6 Luna, global	$9.40	6.3%	2.7%	$94 (9.9%)

Model setup	Per customer-month	Share of Starter $149	Share of Growth $349	Agency, 10 accounts (share of $949)
Gemini 3.8 Flash with 3.5 Flash-Lite, prices from 1 Jan 2027	$7.92	5.3%	2.3%	$79 (8.3%)
Gemini 3.5 Flash, Canada-only on reserved capacity (estimate, about 100 customers)	about $23–31	16–21%	7–9%	$235–312 (25–33%)
Opus 5.5 for every job	$20.52	13.8%	5.9%	$205 (21.6%)

Monthly AI cost per customer, by model setup



US dollars, on the workload above: 30 conversations a month, daily recommendations, drafts and routing. The Canada-only bar is a range because reserved capacity is a fixed monthly cost shared by roughly 100 customers.

Cost is not what limits the choice. What could break the budget is how the system is run, not which model it uses. Without prompt caching, Sonnet chat rises from $8.14 to **$26.61** (before the 10% US-region charge). At high reasoning effort it rises to $14.44. A customer who chats three times as much costs about $30 in the US region, 20% of the Starter price[11]. Caching the fixed part of the prompt (Blueprint §7.3), low or medium effort by default, and the per-customer token caps already in the Blueprint keep the bill near the table. Gemini 3.8 Flash costs $0.75/$3.75 per million tokens on Google's global endpoint until 31 December 2026 ($0.825/$4.125 on regional endpoints), and twice that from 1 January 2027[19]. Its cheapness is temporary, and it offers no Canadian processing either.

GPT-6.1 Sol is the serious alternative, and the case for it is real. It costs the same per token. It scored higher on Artificial Analysis's index at high effort (50 against 47) and used fewer tokens per task ($0.32 against $0.88)[13]. The earlier GPT-6 Sol had the lowest summary-hallucination rate of any current flagship on Vectara's board, 6.5%, against 9.8–12.2% for Claude 4.x, and no Claude 5.x model has been scored there

yet[20]. Its strict output mode also accepts a `pattern` rule that could hold headlines to 30 characters as the text is written[21]. Against it: its first reply is more than four times slower, it is weaker on multi-turn tool use on the last independent board, it is not offered on Google Cloud, and moving to it would mean rewriting the consent text and the list of companies that process data. The way to settle this is a one- to two-day comparison before committing. It should use about 150 real Otto questions, run against a frozen copy of the Google Ads test accounts (never the database with demo or real sign-in data), and score the right tool and date range, invented numbers, ad-text length pass rate, p50 and p95 latency, and cost per conversation. Gemini 3.5 Flash in Montréal belongs in that run, so the client can see what a Canada-only Otto costs in answer quality. A thin internal layer between PPCWay and the model provider (Blueprint §7.6) keeps a later switch to a contained piece of work.

Two details decide how well Otto handles 12 languages and Google's character limits. Smaller models lose more quality outside English. On Anthropic's own translated tests, run with extended thinking on the previous generation, Haiku 4.5 kept 95.7% of its English score in French, 94.3% in German and 92.4% in Hindi, while Sonnet 4.5 kept 97.0–98.2% in the major European languages[22]. Anthropic publishes no per-language results for the 5.5 models. So customer-facing text in any language other than English should come from Sonnet, with Haiku kept to internal routing. The reply language should be set in the system prompt, and each language should be tested with about 50 Otto questions, starting with Canadian French. On character limits, Anthropic's structured outputs guarantee the JSON shape but not string length, and they only accept `minItems` of 0 or 1[23]. Google's limits are 30 characters per headline and 90 per description[24], with 3 to 15 headlines and 2 to 4 descriptions per ad[25]. Double-width characters, as in Chinese, Japanese or Korean, count as two[26]. The Blueprint's checks written in code (§7.4) stay the real gate. They allow two rewrites with the specific fault named, then drop the asset, then fall back to a human-written template.

3

No purpose-built Google Ads model beats a general model with good tools

Nothing on the market built for Google Ads should replace a general model, but two Google tools are worth a look alongside one.

Google offers three kinds of AI help for Ads, and none of them can serve as Otto. Its in-product assistant, Ads Advisor, runs on "the latest Gemini models". In November 2025 Google said it would reach all English-language accounts in early December[27,28]. Its May 2026 successor, Ask Advisor, is "currently available in beta for English language accounts", and Google has announced no API or access for outside developers[29]. Google's open-source developer tools are themselves a general model plus tools. The official Google Ads MCP server is read-only and Apache-licensed. It offers a `search` tool that builds GAQL queries from structured fields, plus field metadata[30]. The Developer Assistant v4 now runs inside Claude Code. It grounds itself in Google's API definitions rather than a special model, and it saves code that would change an account without ever running it[31]. For turning questions into GAQL, Hugging Face has only small unofficial datasets and a few community fine-tunes, none with a model card or published results. No benchmark or paper exists[32].

Product or tool	What it is	Verdict for Otto
Ads API `GenerateText`	Google's own generator of headlines, long headlines and descriptions for Search, Performance Max, Display and Demand Gen, from a landing-page address, keywords or a prompt of up to 1,500 characters. In **closed beta** since API v22[33,34]	**Apply for the beta.** If admitted, call it only from inside the gateway (the one door to Google) as a second source of drafts, put its output through the same checks, and compare blind against Opus drafts
Google Ads API Developer Assistant v4	Coding plugin for Claude Code and Antigravity	**Use it** while building the sync and gateway code; it has no part in the running app
googleads/google-ads-mcp	Read-only GAQL server	Reference only; at runtime it would reach Google directly and break the one-door rule
Ads Advisor / Ask Advisor	Gemini assistant inside Google Ads, free, English only, beta	A competitor to beat on plain language, approvals, guardrails, audit trail and 11 more languages
Conversational Analytics API	Google-hosted "ask your data" agent, generally available since June 2026 for BigQuery and Looker; Cloud SQL for PostgreSQL is still in preview[35]	Not for v1. It ties Otto to Gemini, gives PPCWay less control over prompts, and the mirror is on Neon, not Cloud SQL

Product or tool	What it is	Verdict for Otto
Arctic-Text2SQL-R1, OmniSQL, SQLCoder	Open models trained to write SQL	Skip. Arctic's 32B led BIRD's single-model board at 71.83% in May 2025[36]. The 2026 leaders are multi-step agent systems built on general models, such as the 76.23% top entry on the harder Spider 2.0-Lite test[37]
Persado, Jacquard (Phrasee)	Enterprise copy optimisation, mainly email and SMS	Skip. Both sell on enterprise quotes[38], with no evidence on 30-character search headlines

The competitors confirm this. Opteo's chat lets an "Orchestrator" choose the model automatically, names OpenAI, Anthropic and Gemini among its providers, and pulls account data with tool calls. A change "goes live immediately" once the user confirms[39]. Optmyzr's Sidekick does not say which model it uses[40]. None of the PPC tools examined has its own ads model. PPCWay's mix of approvals, guardrails and an undo window is stricter than Opteo's confirm-and-apply, and that difference is worth saying plainly in Otto's own copy. Several competitors now offer MCP servers so their tools can be driven from Claude or ChatGPT. PPCWay should not host one that re-exposes Google Ads to outside tools. Google's API policies bar a hosted MCP server or proxy that "solely replicates, wraps, or re-exposes Google Ads programmatic capabilities" unless Google has reviewed it[7]. The evidence on AI against human search ad copy is thin and mixed. One small 8-week test from April 2024 ran one ad of each on a $500 budget. It found a 4.98% click-through rate for human copy against 3.65% for AI copy, on older models[41]. What improves ad drafts is the process around the model: grounding in the business's own site, character and policy checks, and feeding live asset results back in.

Fine-tuning a model on Google Ads expertise is not worth doing now. For adding knowledge, one study found that "while unsupervised fine-tuning offers some improvement, RAG consistently outperforms it"[42]. Google Ads facts also change too fast for trained-in knowledge to stay correct. Google shipped four major API versions, v22 to v25, between October 2025 and July 2026[34]. Otto should instead draw on a dated, curated library of Google help and policy pages and PPCWay's own playbooks. The guardrail rules belong in the instructions, not in the model's training. Fine-tuning becomes worth a look later, as a way to cut cost on the nightly batch once thousands of approved and rejected examples exist. One item belongs on the launch list. In July 2026 Google added a way for advertisers to declare ads made with outside AI tools[43]. The API has accepted that declaration since v25, through the `synthetic_content_info` field on ads and assets[34], so ad text Otto drafts should carry it from the start.

4

Only one ageing Gemini model keeps Otto's processing inside Canada

No current frontier model can be called with a promise that the processing happens in Canada, on any cloud.

On Vertex AI, which Google now calls the Gemini Enterprise Agent Platform, Toronto has no generative AI endpoint at all. In Montréal, only **Gemini 3.5 Flash, Gemini 2.5 Flash and Gemini 2.5 Pro** carry a Canadian processing tick for text, besides two older embedding models[44]. The two 2.5 models retire on **20 October 2026**, six days from now, which leaves 3.5 Flash as the only one[45]. Every newer Gemini (3.6, 3.7 and 3.8 Flash and 4 Argon) skipped Canada. No Claude, Mistral or open model has a Montréal tick: neither the partner-model nor the open-model residency table has a Canada column[46]. The other routes are no better. On AWS Bedrock, Claude 5.5 is reachable from Canada only through a US cross-region profile that sends requests to US regions, or through Global[47]. Azure has no Canada data zone, and its pay-per-use models processed in Canada are only gpt-4o and gpt-4.1-mini[48]. Anthropic's own API offers "us" or "global" only[49]. OpenAI can store data at rest in Canada, but its Canadian "regional processing" is "None"[50].

Gemini 3.5 Flash is a Canada-only option on paper, but in practice live chat cannot simply be bought there. Google's capabilities page lists only Batch and Single Zone Provisioned Throughput as guaranteed in-region ways to use a regional endpoint outside the US and EU, and marks standard pay-per-use "Not supported"[51]. The model's own card lists pay-per-use for the global, US and EU endpoints only. Its list of regions for reserved capacity leaves out Montréal[52]. So the overnight recommendation batch can run in Canada, at $0.825/$4.95 per million tokens. Live chat would need a reserved single-zone capacity unit, sold only through a Google account representative, and Google must first confirm it is offered for this model in Montréal[53]. One unit costs about **$2,200–$2,970 a month** before any traffic, on a one-year or one-month commitment[18]. It handles 675 tokens a second, with a cached input token counting as 0.1 and an output token as 6[54]. On the chat workload above, with peaks about four times the average, one unit covers roughly 100 customers. That puts a Canada-only Otto at **about $23–31 per customer a month**, more than twice the recommended setup. This is an estimate; the real figure depends on how bursty chat traffic turns out to be. The model is also behind the current lineup. It scores 33 on Artificial Analysis's current index at high effort, below Sonnet 5.5 at low effort (36). It takes about 17 seconds to start replying at high effort, and Artificial Analysis now labels it deprecated[55]. Google supports it until at least **19 May 2027**. The fair case for it is that it handles tools well. Google reports 83.6% on the MCP Atlas tool-use test, above the Claude 4.x and GPT-5.5 models it was compared with[56], and function calling and structured output are both guaranteed in region. Because accuracy comes mostly

from the tools, a Canada-only Otto on 3.5 Flash would answer questions acceptably. It would cost more, be slower, write weaker drafts and explanations, and face a forced migration in about seven months.

Running an open model on PPCWay's own servers is the only other Canada-only route. Toronto offers L4 and H100 GPUs for Vertex model hosting[57]. A single always-on L4 machine lists at about $593 a month at US-central rates[58], so one or two machines come to roughly $600–$1,700 a month before any backup capacity. That is a fixed cost large next to a per-customer bill. It would also bring a weaker model and server upkeep the owner has so far avoided.

5

The client's decision is what "stored in Canada" has to mean

The landing page promises that data is stored in Canada, and the client has to decide whether that promise also covers where Otto's AI processes it.

The requirements document leaves this open as "one region, or Canada, US and EU regions?" (Requirements). The product already sets a pattern for it. The designed settings screen says data lives in Canada, "Montreal and Toronto", and lists Anthropic as "Outside Canada" for "site text only" (settings screen). Otto will send Google Ads data to the model: campaigns, search terms and results. Whichever option is chosen, if it is not Canada-only, that row has to change.

The law allows US processing, provided PPCWay says so openly. Canada's privacy commissioner treats sending data to a processor abroad as a "use", not a disclosure, so no extra consent is needed. Organisations must still "make it plain" that information "may be processed in a foreign country" and may be reachable by that country's authorities[59]. Quebec's Law 25 needs a privacy impact assessment and a written agreement before personal information leaves Quebec, even for another province[60]. Because PPCWay's servers are in Toronto, that assessment is needed for Quebec customers under every option. Montréal processing alone would not trigger it, since Montréal is in Quebec. The Google Ads API Terms say nothing about AI. The API policies do require apps to tell users in the app when data serves a "secondary purpose", and agencies need written client consent before "otherwise disclosing" account data[7]. The privacy commissioner's AI principles, and the EU AI Act's Article 50 since 2 August 2026 (not delayed by the 2026 amending regulation), both expect users to be told they are talking to an AI[61,62].

	A. Stored in Canada, AI processed in the US (recommended)	**B. Everything in Canada**	**C. Each customer in their own region**
Models	Claude Sonnet, Haiku and Opus 5.5 on Vertex `us`	Gemini 3.5 Flash in Montréal only, or a self-hosted open model in Toronto	Claude on Vertex `us` and `eu`; Canadian customers as in A or B
Quality and future	Current frontier models, with OpenAI and Google as fallbacks	One deprecated-labelled model, supported to May 2027, no successor announced	As A
LLM cost per customer-month	About $11.60	About $23–31 (estimate), from reserved capacity at $2,200–2,970 a month per unit	About $11.60

	A. Stored in Canada, AI processed in the US (recommended)	B. Everything in Canada	C. Each customer in their own region
What the client must accept	Reworded promise and a US processor in the privacy notice	A reserved-capacity contract that Google must first confirm it sells in Montréal; slower chat, weaker drafts and a migration by mid-2027	Separate regional databases and EU (GDPR) obligations

> **Recommendation: choose Option A, and change the wording rather than the model.** The promise would read: "Your data is stored in Canada. When you ask Otto something, the figures it needs are sent to Google Cloud in the United States to write the answer, and are not kept there or used for training." That wording extends the "Outside Canada" pattern the product already uses. It costs less than half as much per customer as Option B, needs no capacity contract, and avoids tying Otto to a model already labelled deprecated. Claude should be called through Vertex `us` rather than Anthropic's own API. Both cost 10% over global routing[11]. On Vertex, though, Google is the data processor[63], and Google Cloud already hosts PPCWay. Usage is billed to the client's own Google Cloud project (decision record). Google commits Haiku, Sonnet and Opus 5.5 to the US and EU multi-regions, which also covers Option C later[44]. Anthropic's own pages, by contrast, give conflicting answers on how long API data is kept: one says it is not kept by default, the other that it is deleted within 30 days[63,64].

Option A needs six things done before launch to keep the promise true:

1. Block the global endpoint with an organisation policy, since data sent there "may be processed in any Google Cloud location around the world"[44].
2. Leave Vertex's request-response logging switched off. It is off by default, and when on it writes to PPCWay's own BigQuery table[65].
3. Ask the Google account team in writing whether zero data retention is possible for Claude. Google's page says that for some advanced models it "may not be possible" and to contact the account team[65].
4. Meet the requirement for prompt logging and policy filtering with two pieces kept in Canada:
   - Model Armor in Toronto. It screens prompts and replies for injection, sensitive data and harmful content, and stores nothing unless logging is switched on. Toronto is a "limited-support" region: it has no malicious-link filter and no multi-language detection, so French prompts need their own test[66,67,68].
   - PPCWay's own redacted, append-only prompt log.
5. Confirm that the Neon database sits in a Canadian region. The decision records do not say where it is, and the mirror and the prompt log will both live there.
6. Update the privacy notice and the processor list, and complete the Quebec assessment.

Option B is right only if winning customers depends on “processed in Canada”, for example regulated or public-sector buyers. In that case, first get Google’s written confirmation that reserved capacity for Gemini 3.5 Flash is sold in Montréal and at what price, and plan the move off 3.5 Flash now.

Conclusion

Strict Canadian residency costs more, not less.

A Canada-only Otto needs reserved capacity, roughly $23–31 per customer at launch scale against about $11.60, and Google has not confirmed it sells that capacity in Montréal for the one eligible model. It also ties Otto to a model whose support may end after May 2027, from a vendor whose last four releases all skipped Canada. That makes it a decision about how PPCWay presents itself, and one with a real price. The comparison test should run before the client decides, with Gemini 3.5 Flash included, so the trade-off is put in terms of answers Otto gets wrong rather than in the abstract. If Google later offers a newer Gemini or Claude in Montréal on pay-per-use terms, Option B becomes cheap to adopt, and the thin layer between PPCWay and the model provider is what makes that switch a small job.

The longer-term lesson is that Otto’s lasting value lies in what surrounds the model. The typed tools, the number check, the guardrails, the approvals and the evaluation set built from test accounts all survive a model change. A model choice lasts only until the next release, and all three labs replaced their whole lineup in five weeks this autumn. PPCWay should therefore treat model choice as a quarterly re-run of the same test, not a one-time decision. Its position against Google’s free, English-only Ask Advisor should rest on what Google does not offer: plain-language reasons, approval before anything changes, undo, and 12 languages.

What the fact-check changed

The report's claims were checked against their live sources on 14 October 2026.

The recommendation, Claude on Vertex AI with storage kept in Canada, did not change. These points did:

1. **Canada-only costs more, not less.** Gemini 3.5 Flash cannot be bought pay-per-use in Montréal, so live chat there needs reserved capacity: an estimated $23–31 per customer, not about $8. Google has not confirmed that capacity is sold in Montréal for this model.
2. **Nightly batch on Google Cloud uses Google's own batch prediction**, not Anthropic's Message Batches API, which Google Cloud does not offer. The 50% discount and the cost totals are unchanged.
3. **Google's `GenerateText` is in closed beta.** PPCWay must apply before it can try it.
4. **Google's API policies were quoted too broadly.** MCP servers and automated changes are barred only without Google's review or per-account sign-in, and "immutable transaction logs" is a test in that review, not a blanket rule.
5. **Latency figures were updated** to Artificial Analysis's 14 October measurements. The ranking did not change.
6. **Five quotations were corrected** to the sources' exact wording: dbt, Tableau Pulse, proof-carrying numbers, Beurer-Kellner and Ovadia. A sixth, the data-residency quote in step 1, now carries Google's full wording.
7. **Smaller corrections.** The heavy-chatter cost is about $30, not $27. Ad-copy vendor prices are no longer published. Small GAQL datasets do exist on Hugging Face. The Spider 2.0 and Arctic benchmark claims were narrowed. Model Armor's Toronto limits were added.

The source checks are saved next to the research notes in `research_notes/LLM choice for Otto AI chatbot/`, in the files starting `factcheck_`.

Sources

Numbers in the text point here.

1. **Dbt Labs**
   docs.getdbt.com/blog/semantic-layer-vs-text-to-sql-2026
2. **Rumiantsau & Fokeev**
   arxiv.org/abs/2604.25149
3. **BIRD**
   bird-bench.github.io
4. **Shopify Engineering**
   shopify.engineering/building-production-ready-agentic-systems
5. **Salesforce**
   salesforce.com/blog/tableau-pulse
6. **Solatorio**
   arxiv.org/abs/2509.06902
7. **Google Ads API policies**
   support.google.com/adspolicy/answer/6169371
8. **OWASP**
   genai.owasp.org/llmrisk/llm01-prompt-injection
9. **Beurer-Kellner et al.**
   arxiv.org/html/2506.08837v1
10. **Anthropic models overview**
    platform.claude.com/docs/en/about-claude/models/overview
11. **Anthropic pricing**
    platform.claude.com/docs/en/about-claude/pricing
12. **Anthropic**
    anthropic.com/news/claude-haiku-5-5
13. **Artificial Analysis**
    artificialanalysis.ai/leaderboards/models
14. **BFCL**
    gorilla.cs.berkeley.edu/data_overall.csv
15. **Anthropic**
    anthropic.com/news/claude-sonnet-5-5
16. **Claude on Vertex AI**
    platform.claude.com/docs/en/build-with-claude/claude-on-vertex-ai
17. **Claude batch on Google Cloud**
    docs.cloud.google.com/gemini-enterprise-agent-platform/models/partner-models/claude/batch
18. **Pricing**
    cloud.google.com/gemini-enterprise-agent-platform/generative-ai/pricing
19. **Vertex pricing**
    cloud.google.com/vertex-ai/generative-ai/pricing
20. **Vectara**
    github.com/vectara/hallucination-leaderboard
21. **OpenAI**
    developers.openai.com/api/docs/guides/structured-outputs
22. **Anthropic**
    platform.claude.com/docs/en/build-with-claude/multilingual-support
23. **Anthropic**
    platform.claude.com/docs/en/build-with-claude/structured-outputs
24. **System limits**
    developers.google.com/google-ads/api/docs/best-practices/system-limits
25. **Responsive search ads**
    developers.google.com/google-ads/api/docs/responsive-search-ads/overview
26. **Google Ads Help**
    support.google.com/google-ads/answer/7684791
27. **Google**
    blog.google/products/ads-commerce/ads-advisor-and-analytics-advisor
28. **PPC Land**
    ppc.land/google-launches-ads-and-analytics-advisors-for-all-english-accounts
29. **Google**
    blog.google/products/ads-commerce/ask-advisor
30. **GitHub**
    github.com/googleads/google-ads-mcp
31. **Google Ads API docs**
    developers.google.com/google-ads/api/docs/developer-toolkit/ai-assistant
32. **Hugging Face**
    huggingface.co/models?search=gaql
33. **Google Ads API**
    developers.google.com/google-ads/api/reference/rpc/v25/GenerateTextRequest
34. **Release notes**
    developers.google.com/google-ads/api/docs/release-notes
35. **Release notes**
    docs.cloud.google.com/gemini/data-agents/conversational-analytics-api/release-notes
36. **Snowflake**
    snowflake.com/en/blog/engineering/arctic-text2sql-r1-sql-generation-benchmark
37. **Spider 2.0**
    spider2-sql.github.io
38. **AI Agent Index**
    theaiagentindex.com/compare/jacquard-vs-persado
39. **Opteo**
    opteo.com/docs/en/articles/12968321-ai-chat
40. **Optmyzr**
    help.optmyzr.com/en/articles/8743581-optmyzr-sidekick-ai-faqs
41. **Search Engine Journal**
    searchenginejournal.com/humans-vs-machines-ad-copy-content-test-data-study/509942
42. **Ovadia et al.**
    arxiv.org/abs/2312.05934

43. **Google**
blog.google/products/ads-commerce/google-ads-ai-transparency-labels

44. **Data residency**
docs.cloud.google.com/gemini-enterprise-agent-platform/resources/data-residency

45. **Model versions**
docs.cloud.google.com/gemini-enterprise-agent-platform/models/model-versions

46. **Locations**
docs.cloud.google.com/gemini-enterprise-agent-platform/resources/locations

47. **AWS**
docs.aws.amazon.com/bedrock/latest/userguide/model-card-anthropic-claude-sonnet-5-5.html

48. **Microsoft**
learn.microsoft.com/en-us/azure/foundry/foundry-models/concepts/models-sold-directly-by-azure-region-availability

49. **Anthropic**
platform.claude.com/docs/en/manage-claude/data-residency

50. **OpenAI**
developers.openai.com/api/docs/guides/your-data

51. **Supported capabilities**
docs.cloud.google.com/gemini-enterprise-agent-platform/resources/supported-capabilities

52. **Gemini 3.5 Flash card**
docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-5-flash

53. **SZPT**
docs.cloud.google.com/gemini-enterprise-agent-platform/models/provisioned-throughput/szpt

54. **Provisioned Throughput**
docs.cloud.google.com/gemini-enterprise-agent-platform/models/provisioned-throughput/supported-models

55. **Artificial Analysis**
artificialanalysis.ai/models/gemini-3-5-flash

56. **DeepMind**
deepmind.google/models/model-cards/gemini-3-5-flash

57. **ML locations**
docs.cloud.google.com/gemini-enterprise-agent-platform/machine-learning/general/locations

58. **Pricing**
cloud.google.com/products/gemini-enterprise-agent-platform/pricing

59. **OPC**
priv.gc.ca/en/privacy-topics/airports-and-borders/gl_dab_090127

60. **McCarthy Tétrault**
mccarthy.ca/en/insights/blogs/techlex/quebec-privacy-regulators-guide-law-25-compliant-privacy-impact-assessments

61. **OPC**
priv.gc.ca/en/privacy-topics/technology/artificial-intelligence/gd_principles_ai

62. **AI Act**
artificialintelligenceact.eu/article/50

63. **Anthropic**
platform.claude.com/docs/en/manage-claude/api-and-data-retention

64. **Anthropic Privacy Center**
privacy.claude.com/en/articles/7996866-how-long-do-you-store-my-organization-s-data

65. **Zero data retention**
docs.cloud.google.com/gemini-enterprise-agent-platform/resources/zero-data-retention

66. **Model Armor locations**
docs.cloud.google.com/model-armor/locations

67. **Features by region**
docs.cloud.google.com/model-armor/feature-availability-by-region

68. **Model Armor**
docs.cloud.google.com/model-armor/overview